

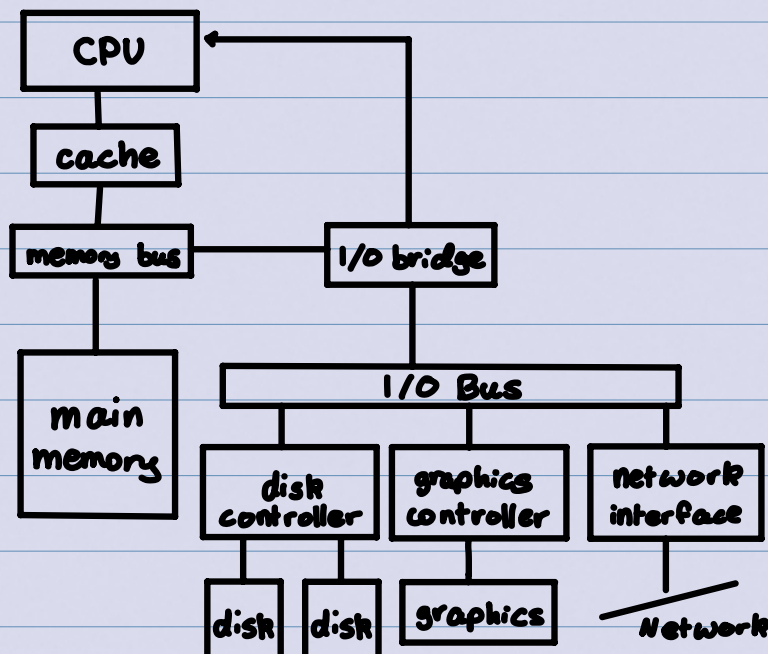
# Lecture 1 - 4<sup>th</sup> Sept 2024

## Instruction Set Architecture: how hardware executes software

### • What is a computer?

- ↳ just a machine that does whatever the software tells it to do.
- ↳ computer > digital circuits > logic gates > transistors
- ↳ software is just a series of instructions.
- ↳ The five classic components of a computer:
  - CPU: Central Arithmetic (ALU) and Central Control
  - Memory System
  - Input and output

↳ But, here's a more realistic view:



# Architecture vs Microarchitecture

## ↳ Processor architecture:

- Functional appearance to software (ISA)
  - Exactly what instructions does it have?
  - Number of memory/storage locations it has
  - Interface!

in this case, compilers!


## ↳ Processor microarchitecture:

- Logical Structure that implements the architecture
  - Number of functional units, interconnection, control
  - Size of the caches
  - Not visible to the software
  - Implementation!

## Memory Hierarchy:

- CPU
- CPU Cache
- Physical Memory → RAM
- Solid State Memory → SSD/USB stick
- Virtual Memory → hard drive

processor




## • Moore's Law: every 18-24 months,

- 2x transistors on same chip area
- 2x processor speed
- 2x memory capacity
- $\frac{1}{2}$  energy/power consumption.

↳ technically not a law, but it's a self-fulfilling prophecy.

Memory Wall: every 2 years,

- 2x Instructions / second
- 2x memory capacity
- 1.1x memory latency



↳ growing disparity between processor and memory performance!  
∴ significant effort in reducing/hiding memory latency.

Latency: time from start to finish (aka response time).

Throughput: number of tasks completed per time unit  
(aka bandwidth)

↳ throughput can exploit parallelism, but latency cannot!

- Improving the latency of a component always improves the overall system throughput.

Speedup =  $\frac{\text{Performance}(x)}{\text{Performance}(y)}$  → "x is 'speedup' times faster than y"

↳ can be broken down into  $\frac{\text{throughput}(x)}{\text{throughput}(y)}$  or  $\frac{\text{latency}(y)}{\text{latency}(x)}$ .  
don't forget - latency will likely be < 1, so need to invert fraction!

Iron Law of Performance:

CPU Time:  $\frac{\text{instructions}}{\text{program}} \cdot \frac{\text{cycles}}{\text{instruction}} \cdot \frac{\text{seconds}}{\text{cycle}}$   
aka CPI

IPC is inverse - instructions per cycle